

Advanced OpenMP Features: Tasking and Accelerator Support

Piotr Luszczek

OpenMP Tasking Overview

- OpenMP with version 4 started supporting tasking with data dependences
- Tasks were part of OpenMP since version 3
- Version 4 introduced data dependence clause
 - New clause “depend” allows to constraint execution of tasks based on how they pass data between each other
 - The data dependence also allows the tasks to synchronize selectively between each other based on the data they exchange
- OpenMP provides a scheduler that decides the cores to run the tasks
 - The scheduling decisions are made based on
 - Amount of work each core has
 - Availability of data in caches
 - It is beneficial to use the same core for tasks that share data that is already in cache

Task Directive Summary

- `#pragma omp task`
 - `if (expression)`
 - `final`
 - `untied`
 - `default`
 - `mergeable`
 - `private (list)`
 - `firstprivate (list)`
 - `shared (list)`
 - `depend (dep. type : list)`
 - `priority (expression)`
- By far, task might be the most complicated out of the commonly used ones
- The most important clauses will be discussed next

Depend Clause Summary

- `#pragma omp task depend(dependence type : list)`
- Dependence type can be either
 - `in`
 - Use it for data that is consumed by the tasks
 - `out`
 - Use it for data that is produced by the tasks
 - `inout`
 - Use it for data that is both consumed and produced by the tasks
 - Recall Fortran specification for subroutine function parameters
 - Since Fortran 90: `in`, `out`, `inout`, `scratch`
- List is a comma-separated enumeration of variables participating in data dependence graph
 - Arrays are specified with ranges:
 - For example: `depend(in:A[0:N])`

Recall DAG of Tasks and Its Level Sets

Level Set 1

Load(A3)

Load(A1)

Load(A2)

Level Set 2

Add(A1, A2) A4

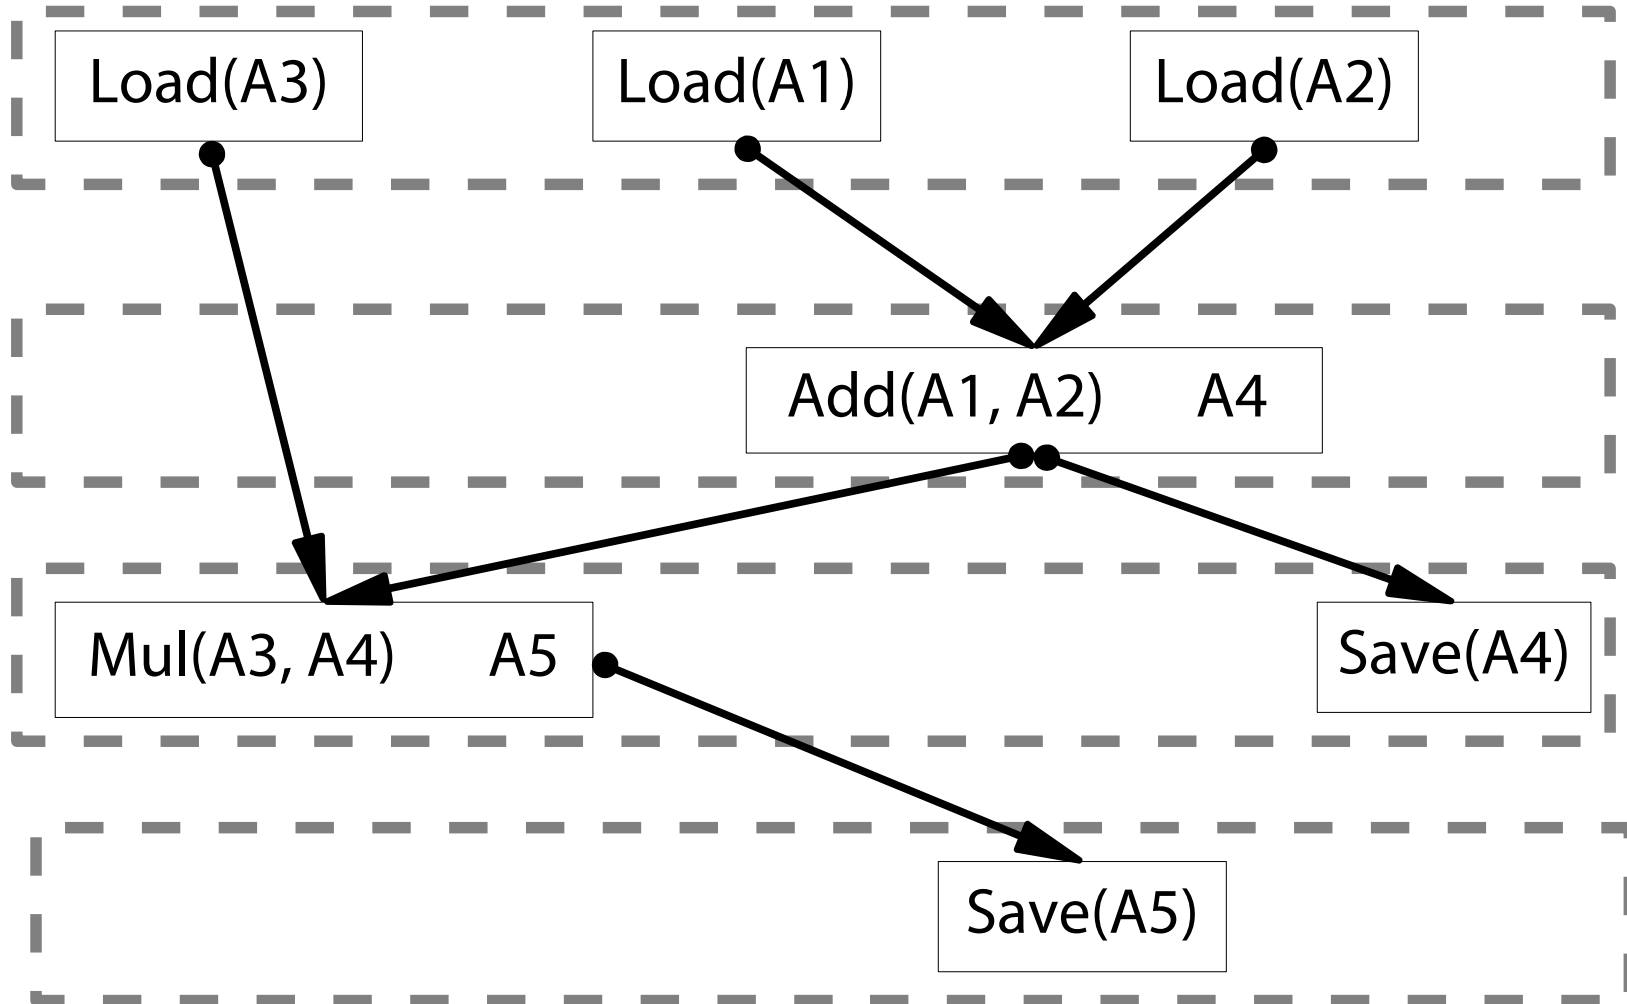
Level Set 3

Mul(A3, A4) A5

Save(A4)

Level Set 4

Save(A5)



Implementation of the Sample Task DAG

```
#pragma omp parallel shared(a1, a2, a3, a4, n)
{
    #pragma omp master
    {
        #pragma omp task depend(out:a1[0:n])
        Load(n, a1);
        #pragma omp task depend(out:a2[0:n])
        Load(n, a2);
        #pragma omp task depend(out:a3[0:n])
        Load(n, a3);
        #pragma omp task depend(in:a1[0:n],a2[0:n]) depend(out:a4[0:n])
        Add(n, a1, a2, a4);
        #pragma omp task depend(in:a3[0:n],a4[0:n]) depend(out:a5[0:n])
        Mul(n, a3, a4, a5);
        #pragma omp task depend(in:a4[0:n])
        Save(n, a4);
        #pragma omp task depend(in:a5[0:n])
        Save(n, a5);
    }
}
```

- The order in which tasks are inserted is important
 - To be safe, exhaust the tasks from the first level set before inserting tasks from the second level

Advanced Use of Tasks

- It is possible to make complicated task graphs
 - `taskwait` directive creates a join point of descendent tasks of the current task
 - `taskgroup` allows creating of and waiting for descendant tasks
 - `taskyield` directive allows to relinquish CPU for other tasks
 - `taskloop` directive allows generation of tasks with loop constructs

OpenMP 4+ is not the Only One for Tasking

- Task-based data-parallel computing has become an important method for parallel computing
 - It has grown tremendously since the multicore era began
- There are many projects that use this paradigm
 - Shared memory: Hstreams, Open Community Runtime, ompSS, QUARK, Thread Building Blocks
 - Distributed memory: Legion, PaRSEC, RAJA, StarPU, Thor
 - There are more projects that provide this functionality

OpenMP and Accelerators

- It is possible to offload OpenMP tasks to accelerators
 - “target data” directive maps variables to a device
 - `#pragma omp target data if(scalar expression) device(integer expression) map(map type: list) use_device_ptr(list)`
 - “target enter data” performs the mapping and “target exit data” unmaps them
 - “target update” assures data consistence between host and device
 - “target” directive can be used to map data and execute code on device
- Support for OpenMP for accelerators is limited due to the market situation
 - GNU Compiler Collection include support for OpenMP offload but it usually is disabled and cannot be used at the same as OpenACC
 - Intel discontinued support for Xeon Phi accelerators beyond Knights Landing and Knights Hill hardware will not be released

OpenMP and OpenACC

- OpenMP started as CPU-only specification that now includes accelerators
 - Continues to have support from major CPU vendors
- OpenACC unified many offloading solutions for accelerators and now includes CPU as a target
 - It may be used with GNU and PGI compilers
 - US' DOE sponsors further development
 - This may introduce LLVM support